Chemical Identity Lost in Regulation: A Study of Semantic Interoperability in European Chemical Substance Data

Geert VAN HAUTE [a,1], Pieter FANNES [a] and Stijn GOEDERTIER [a]

[a] *Department of Environment and Spatial Development, Flemish Government, 1000 Brussels, Belgium*

ORCiD ID: Geert Van Haute https://orcid.org/0009-0002-9836-8139, Pieter Fannes https://orcid.org/....-....-....-...., Stijn Goedertier https://orcid.org/....-....-....-....

Abstract. Chemical substances are fundamental entities in environmental regulation, yet their representation in regulatory data systems diverges significantly from established practices in chemistry and cheminformatics. While chemical identity can be unambiguously defined using structure-based identifiers such as the International Chemical Identifier (InChI) and InChIKey, European regulatory datasets rely on heterogeneous identifiers including CAS numbers, EC numbers, textual names, and loosely defined group descriptions.

This paper analyses the resulting interoperability challenges across European regulatory frameworks, based on a case study spanning 18 datasets from the European Chemicals Agency (ECHA) and related sources. We show that regulatory "substances" frequently represent heterogeneous semantic entities, including individual molecules, mixtures, analytical parameters, and substance groups, for which formal relationships are not explicitly defined.

Our results demonstrate that (i) a substantial fraction of regulatory entries cannot be linked to a defined chemical structure, (ii) CAS numbers do not provide a reliable unique identifier, and (iii) semantic and embedding-based approaches fail to recover chemical identity for non-structure substances. We further show that resolving these inconsistencies retrospectively leads to combinatorial complexity that is infeasible at scale.

In contrast, structure-based identifiers and ontology-driven approaches, as used in cheminformatics resources such as PubChem and ChEBI, enable scalable integration through linked data principles. We argue that interoperability of chemical substance data must be addressed at the level of legislation and data modelling, by adopting structure-based identifiers, formal ontologies, and machine-readable representations.

Keywords. chemical substance identification, semantic interoperability, regulatory interoperability, REACH, InChIKey, cheminformatics, FAIR data, knowledge graph, European environmental legislation, Interoperable Europe

[1] Corresponding Author: Geert Van Haute, contact details.

1. Introduction

Chemical substances are central entities in environmental, health, and product regulation. Across European legislation, substances are used to define obligations, assess risks, and communicate hazards in frameworks such as the Registration, Evaluation, Authorisation and Restriction of Chemicals (REACH) [1], the Classification, Labelling and Packaging of chemicals (CLP) [2], and sector-specific regulations including pesticides and water quality directives. Effective governance of these substances increasingly depends on the ability to integrate and reuse data across regulatory domains.

In chemistry and cheminformatics, chemical identity is well-defined and operationalised through algorithmically defined, structure-based, representations. Molecular structures can be uniquely encoded using standards such as SMILES notations [3], IUPAC names [4], the International Chemical Identifier (InChI) [5] and its hashed representation, the InChIKey, enabling unambiguous identification, comparison, and linking of substances across databases [6]. Large-scale infrastructures such as PubChem [7] and ChEBI [8] build on these identifiers to provide machine-readable, ontology-based knowledge graphs that support interoperability and automated reasoning.

In contrast, regulatory data systems adopt a fundamentally different approach. Substances are typically identified using heterogeneous and often ambiguous identifiers, including Chemical Abstracts Service (CAS) registry numbers [9,10], European Community numbers (EC numbers) [11], textual names, and informal group descriptions. Moreover, the concept of a *substance* in regulation is not uniform: it may refer to a single molecule, a mixture of unknown or variable composition (UVCB) [12], a group of related compounds (e.g., PFAS), or even an analytical parameter defined as a function over multiple substances. As illustrated in Figure 1, these entities correspond to fundamentally different semantic types.

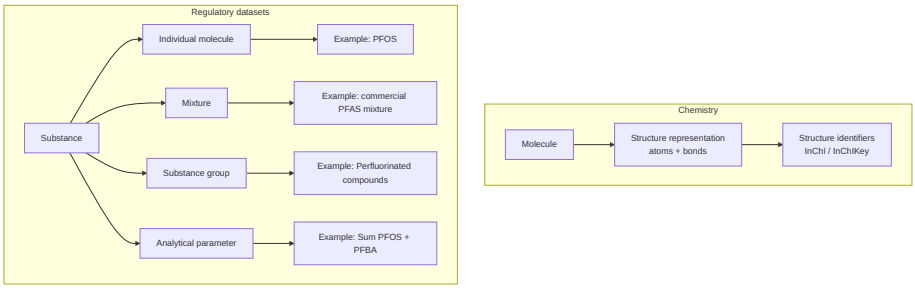

Figure 1. Conceptual model of the semantic mismatch. In chemistry, a substance typically corresponds to a molecule defined by its molecular structure and uniquely identified using structure-derived identifiers such as InChI or InChIKey. In regulatory datasets, however, the concept of a "substance" often encompasses heterogeneous entities including individual molecules, mixtures, substance groups, and analytical sum parameters. These entities correspond to fundamentally different semantic types, which complicates interoperability and automated reasoning across regulatory datasets. This conceptual mismatch is reflected empirically in Section 4, where a large fraction of regulatory entries cannot be linked to a defined molecular structure.

As a result, integrating substance information across regulatory frameworks becomes a complex and largely manual task, hindering policy enforcement for the envi-

ronmental administration (like the Flemish Department of Environment and Spatial Development) and also increases uncertainty for industrial players hindering environmental management, invenstment and innovation.

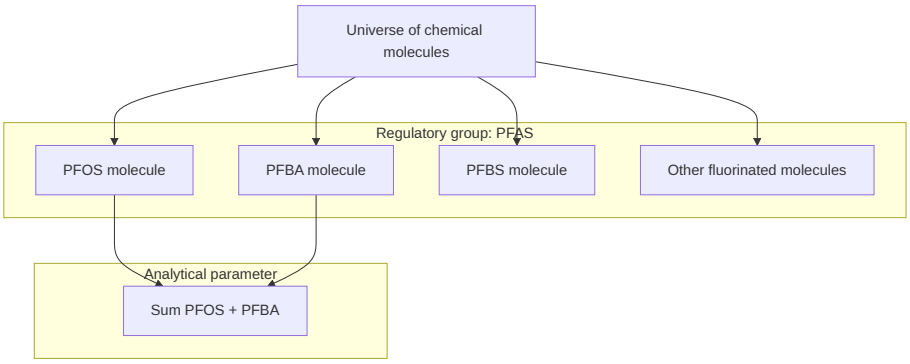

Figure 2. Set-theoretic view of regulatory substances. The lack of explicit set-based semantics leads to unclear relationships between entities and contributes to combinatorial complexity.

This paper investigates the extent and consequences of this interoperability problem using a comprehensive analysis of 18 regulatory datasets from the European Chemicals Agency (ECHA) and related sources. We focus on three main research questions:

- To what extent can regulatory substances be linked to unambiguous chemical structures?
- How consistent and reliable are commonly used identifiers such as CAS numbers for cross-dataset linking?
- Can semantic or data-driven methods compensate for the absence of structure-based identifiers?

Based on this analysis, we make the following contributions:

1. We quantify the heterogeneity of regulatory substance definitions and introduce a taxonomy of linkability reflecting their degree of chemical specificity.
2. We demonstrate the limitations of CAS-based identification and show that a substantial fraction of regulatory substances cannot be resolved to a defined structure.
3. We evaluate semantic embedding approaches and show that they fail to recover chemical identity for non-structure substances.
4. We estimate the scale of the manual effort required for retrospective harmonisation, demonstrating its infeasibility.
5. We show that structure-based identifiers combined with ontology-driven approaches enable scalable integration through knowledge graphs.

The remainder of this paper is structured as follows. Section 3 describes the data sources and analytical methods. Section 4 presents the empirical findings. Section 5 discusses the implications for regulatory data infrastructures and semantic interoperability. Section 6 concludes with recommendations for future work.

2. Related Work

Interoperability of scientific data has been extensively studied in the context of FAIR principles [13] and knowledge graph infrastructures. In the life sciences, large-scale resources such as PubChem [7] and the Chemical Entities of Biological Interest (ChEBI) ontology [8] demonstrate how chemical entities can be represented using structure-based identifiers, ontology-driven classifications, and machine-readable knowledge graphs. Extensive work on rule-based chemical classification has been undertaken by ClassyFire [14]. In this framework, a structure-based taxonomy (ChemOnt) has been established, where each class is characterised by precise and unambiguous formal rules. These infrastructures enable consistent integration, querying, and reuse of chemical information across domains.

Within the Semantic Web community, numerous approaches have been proposed for data integration based on linked data principles, ontology alignment, and entity resolution [14]. Techniques such as identifier mapping, schema matching, and, more recently, embedding-based similarity methods [15] have been applied to address heterogeneity across datasets. However, these approaches typically assume that entities correspond to well-defined and comparable concepts.

In cheminformatics and regulatory science, considerable effort has focused on handling substances without precise structural definitions, such as UVCBs. Methods like those of Kutsarova et al. [12] use generic structural encodings (e.g., G SMILES) and statistical sampling to generate representative constituent structures, reducing combinatorial complexity while preserving structural diversity.

These approaches operate at the level of individual substances and do not address broader semantic inconsistencies across regulatory datasets. Many entries still lack structure-resolvable identifiers or rely on heterogeneous definitions, limiting interoperability between data sources.

In the regulatory domain, previous work has highlighted challenges related to data fragmentation, heterogeneity of identifiers, and limited machine-readability of legal data [16]. In particular, chemical regulation introduces additional complexity due to the use of substance groups, mixtures, and analytical parameters that do not correspond to single chemical structures. While prior studies have focused on improving data accessibility or applying FAIR principles to regulatory datasets, the fundamental mismatch between chemical identity and regulatory substance definitions has received limited attention.

3. Methods

Our analyses comprises of the following steps:

Data Integration Substance records were compiled from a comprehensive set of 18 European regulatory and monitoring sources and one Flemish administrative list. These include ECHA obligation lists (Candidate List [17], Authorisation List [18], Restriction List [19], POPs List [20], EU Positive List [21], and the CLH Harmonised Classification and Labelling List [22]), ECHA activity lists (Restriction Process [23], SVHC Identification [24], Authorisation Process [25], Dossier Evaluation [26], CLH Process [27],

Substance Evaluation [28], POPs Process [29], PBT Assessment [30], and ED Assessment [31]), as well as REACH registration data [32].

These are the most autoritative datasets that lay the foundation of EU regulatory oversight of chemical Substances.

Additional sources comprise the EU Pesticides Database (Active Substances) [33], the OSPAR List of Chemicals for Priority Action [34], and a Flemish administrative substance list developed in the context of water quality modelling [35]. This Flemish `summation substances` list reflects the heterogeneity of *substance names* in the Flemish environmental permits and - monitoring. For each entry, available identifiers (CAS number, EC number, and substance name) were collected.

Where possible, substances were linked to structure-based identifiers (InChIKey) via PubChem [7]. The resulting dataset establishes a unified mapping between regulatory entries, their identifiers, and their originating source lists.

Chemical Identity and Linkability Each substance entry was classified into one of several types reflecting its degree of chemical specificity: structure-defined molecule, substance group, mixture/UVCB, analytical parameter, or administrative entry.

Based on this classification, a linkability taxonomy was defined, ranging from fully structure-resolvable entries (InChIKey) to non-identifiable entries. This taxonomy expresses the extent to which regulatory substances can be unambiguously linked across datasets.

Identifier Consistency Analysis For entries with both CAS numbers and InChIKeys, bidirectional mappings were analysed to assess identifier consistency. This analysis evaluates whether CAS numbers uniquely map to structures and vice versa.

Structure Coverage and Cross-List Overlap For each regulatory source, the proportion of structure-defined substances was calculated. Overlap between sources was analysed using structure-based matching (InChIKey) and set similarity measures.

Analysis of Non-Structure Substances Substances without structure identifiers were analysed using text-based methods, including embedding-based clustering [36,37,38] and similarity matching to chemical taxonomies (ChemOnt), to assess whether semantic techniques can recover meaningful structure or classification [39,14]. Each substance name and ChemOnt class label was encoded using the `all-MiniLM-L6-v2` sentence-transformer model [15,40,41], yielding a 384-dimensional dense vector representation.

Embedding Model Validation Using ChemOnt Ground Truth To evaluate the ability of text-based embeddings to recover chemical classifications from substance names, an internal validation was performed using structure-defined substances as ground truth. Substances with an unambiguous InChIKey were linked to the ChemOnt taxonomy through ClassyFire classification [14], providing reference assignments at the levels of direct parent, subclass, class, superclass, and kingdom.

Substance names and ChemOnt class labels were encoded using pre-trained transformer models (`all-MiniLM-L6-v2, all-mpnet-base-v2, allenai/scibert_scivocab_uncased` and `dmis-lab/biobert-v1.1`) within the `sentence-transformers` framework [15, 40,41]. Cosine similarity was used to rank all ChemOnt classes for each substance name.

Performance was assessed using standard information retrieval metrics, including Hit@k ($k = 1, 3, 5, 10$) and Mean Reciprocal Rank (MRR) [39]. To account for the hierarchical structure of ChemOnt, Hit@1 was additionally computed for each taxonomy

level. A weighted lenient MRR was introduced to summarise performance across levels, assigning ordinal weights of 1.00 (direct parent), 0.75 (subclass), 0.50 (class), 0.25 (superclass), and 0.10 (kingdom), thereby favouring more specific classifications.

This validation provides a principled comparison of embedding models and demonstrates the extent to which semantic representations of substance names approximate structure-based chemical taxonomy assignments.

Knowledge Graph Construction All data were integrated into an RDF knowledge graph using SKOS [42] for taxonomy representation. Structure-defined substances were linked to ChemOnt classes and to ChEBI entities via InChIKey, enabling integration of biological hazard annotations [43].

Workload Estimation For non-structure substances, interoperability requires manual assessment of pairwise relationships (equivalence, subset, overlap, disjoint). The required effort scales quadratically with the number of such entries as $\binom{N}{2} = N(N-1)/2$ where N is the number of non-structure entries under consideration.

Priority Scoring To demonstrate the analytical capabilities of the integrated knowledge graph, a composite priority score was developed by integrating two independent dimensions: (i) regulatory presence across multiple European substance lists, including the CLP Regulation [2], the EU Chemicals Strategy for Sustainability [44], and the Persistent Organic Pollutants (POPs) Regulation [45]; and (ii) biological role annotations obtained from ChEBI [8]. The approach is conceptually grounded in the ToxPi (Toxicological Priority Index) framework [46], which integrates heterogeneous evidence streams into a single prioritisation score through transparent, domain-specific weights.

This score is not intended to constitute a formal risk assessment. Rather, it serves as an illustrative aggregation of heterogeneous regulatory and biological information within a unified semantic framework. The objective is to demonstrate how interoperable, structure-linked data enable cross-domain queries and comparative analyses that are difficult to achieve using fragmented regulatory datasets. Moreover, such an approach can support the prioritisation of policy enforcement activities, particularly in contexts where regulatory agencies operate with limited personnel resources.

Overall, this methodology highlights the value of integrating regulatory metadata with ontology-based annotations to facilitate the identification of substances and chemical classes that are prominent across multiple dimensions of interest.

4. Results

Overview of the integrated knowledge graph We constructed an integrated knowledge graph to analyse the representation and interoperability of regulatory substances across datasets. Figure 3 provides an overview of the resulting graph structure and its main semantic layers.

Dataset Overview Integration of 18 regulatory sources yielded 41,813 records. Of these, 63% could be linked to a defined molecular structure via InChIKey.

Heterogeneity of Substance Definitions A significant share of regulatory entries (37%) does not correspond to a defined chemical structure, consistent with Kutsarova et al. [12]. These include mixtures (UVCBs), substance groups, analytical parameters, or entries

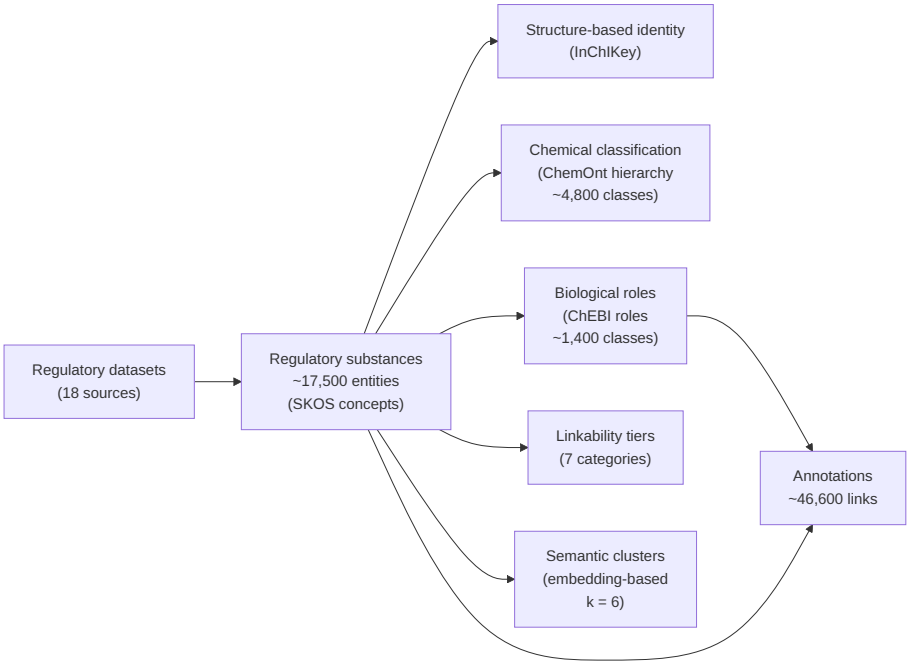

Figure 3. Overview of the integrated knowledge graph constructed from 18 regulatory datasets. Regulatory substances (17,500 entities) are represented as SKOS concepts and form the core layer of the graph. Where possible, substances are linked to structure-based identifiers (InChIKey), enabling alignment with chemical classification systems (ChemOnt) and biological roles (ChEBI). In addition, substances are organised into linkability tiers reflecting their degree of chemical specificity, and into semantic clusters derived from embedding-based similarity. The graph is further enriched with annotations (46,600 links) connecting substances to regulatory sources and functional roles. The figure highlights the multiple semantic layers required to represent and integrate heterogeneous regulatory substance definitions.

without usable identifiers, reflecting differences in how "substances" are defined across regulatory contexts.

Heterogeneity varies widely between regulatory lists: some are mostly structure-based, while others rely heavily on group or mixture definitions, limiting interoperability. (Figure 2)

Applying the seven-tier linkability taxonomy (Section 3) shows that 62.5% of entries (20,909) are directly linkable via InChIKey (tier 1). Only a small fraction consists of partially resolvable groups (tiers 2–3; $<0.2\%$).

The remaining 37.5% cannot be directly linked to a structure: 11.6% have unresolved CAS numbers (tier 4), 10.8% are mixtures or UVCBs (tier 5), $<0.1\%$ are administrative entries (tier 6), and 14.9% lack any usable identifier (tier 7).

Overall, more than one-third of regulatory entries cannot be mapped to a unique chemical structure, and even CAS numbers do not guarantee unambiguous identification. Consequently, structure-based integration applies only to part of the data, requiring alternative approaches for the remainder.

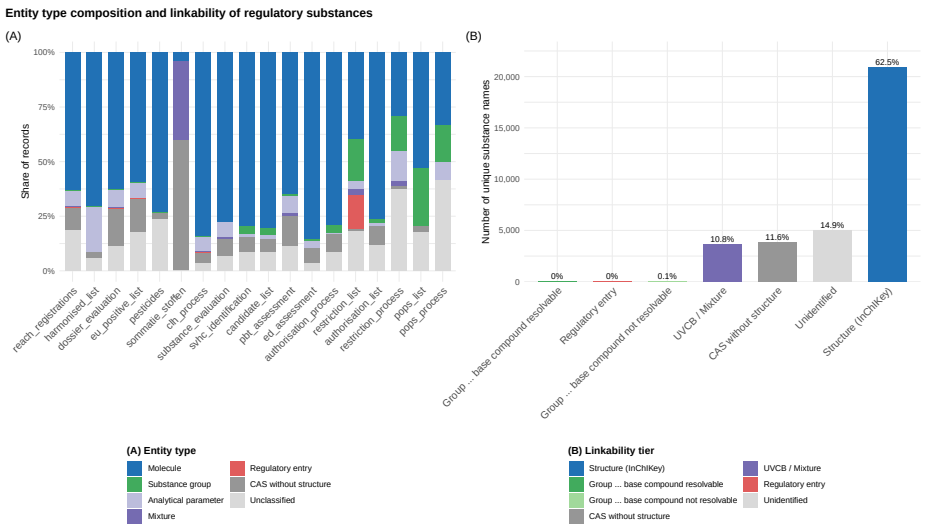

Figure 4. (A) Entity type composition across 18 regulatory sources. Each bar shows the proportional breakdown of substance entries into seven mutually exclusive types (Molecule, Substance group, Analytical parameter, Mixture, Regulatory entry, CAS without structure, Unclassified), illustrating that the mix of entity types differs substantially between regulatory instruments. (B) Linkability taxonomy applied to all 33,452 unique regulatory substance entries. A total of 20,909 entries (62.5%) can be directly linked to a chemical structure via InChIKey (tier 1), while 4,981 entries (14.9%) lack any usable identifier. Together, the panels show that structural linkability is unevenly distributed across sources and that more than one-third of all entries cannot be resolved to a unique chemical structure.

Limitations of CAS Numbers 1

Although CAS numbers are widely used, they do not provide a reliable unique identifier. Multiple CAS numbers may refer to the same structure, and conversely, a single CAS number may correspond to multiple structures. This confirms that CAS-based linking is insufficient for interoperability.

Limited Interoperability Across Regulatory Lists 1

Cross-list analysis shows that most substances are unique to a single regulatory source. Overlap between lists is limited, reflecting the domain-specific design of regulatory instruments rather than a shared chemical reference framework.

Failure of Semantic Approaches for Non-Structure Substances 1 Embedding-based clustering of non-structure substance names yields semantically coherent groups, but does not resolve chemical identity. While clusters reflect lexical or contextual similarity between substance descriptions, they do not correspond to well-defined molecular entities (Figure 5A).

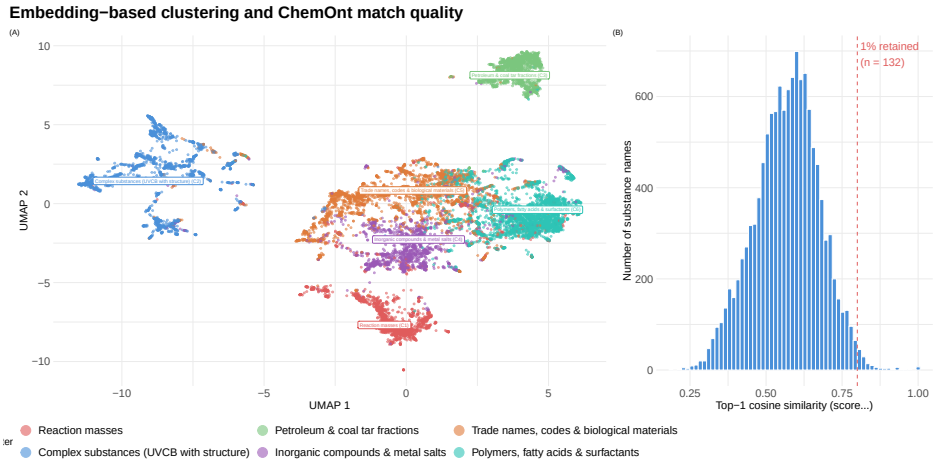

Figure 5. (A) UMAP projection of sentence embeddings for 12,499 non-structure regulatory substance entries, coloured by k-means cluster (k = 6). Each point represents a unique substance name; spatial proximity reflects semantic similarity in the 384-dimensional embedding space. (B) Distribution of top-1 cosine similarity scores [39] between substance name embeddings and ChemOnt class label embeddings across all 12,499 candidate substance-class pairs. The dashed line indicates the applied threshold ($s_1 > 0.80$); 133 entries (1.1%) are retained as matches.

To evaluate whether semantic similarity can recover chemical meaning, we attempted to map non-structure substances to ChemOnt taxonomy classes using cosine similarity between sentence embeddings of substance names and ChemOnt class labels. The distribution of top-1 similarity scores is heavily skewed towards lower values (Figure 5B), with only 133 entries (1.1%) exceeding the threshold $s_1 > 0.80$.

Manual inspection of these 133 retained matches revealed that the majority are chemically incorrect. Correct assignments were limited to cases where the substance name closely resembles the ChemOnt class label (e.g. *Phenols* → *Phenols*, *Alkyl mercury compounds* → *Alkyl mercury compounds*), suggesting that high similarity scores reflect surface-form overlap rather than semantic understanding. In contrast, entries such as *borosilicate glass* → *naphthalene sulfonates*, *polystyrene* → *polysaccharides*, and *silica* → *unsaturated fatty acids* received high scores despite being chemically unrelated. Trade names and reaction products were assigned arbitrary classes with no verifiable basis.

The predominance of incorrect matches above the $s_1 > 0.80$ threshold demonstrates that cosine similarity in embedding space is not a reliable indicator of chemical class membership for non-structure substances. Embedding-based methods cannot compensate for the absence of structure-based identifiers, and fail systematically for substance groups, mixtures, and analytically defined entries.

Validation of Embedding-Based ChemOnt Classification To quantify the ability of text-based embeddings to recover chemical taxonomy assignments from substance names, embedding models were evaluated against a structure-derived ChemOnt ground truth. The validation set comprised 28,204 structure-defined substances with ClassyFire-based assignments to the ChemOnt hierarchy.

Model performance was assessed using standard information retrieval metrics, including Hit@k and Mean Reciprocal Rank (MRR). Among the evaluated models, `all-MiniLM-L6-v2` consistently outperformed the others, achieving a Hit@1 of 5.75%, Hit@3 of 9.86%, Hit@5 of 12.34%, and Hit@10 of 16.21% for the direct parent ChemOnt class. The corresponding MRR was 0.093, indicating that the correct class was typically ranked around the top ten positions. In contrast, `all-mpnet-base-v2`, `allenai/scibert_scivocab_uncased`, and `dmis-lab/biobert-v1.1` showed substantially lower performance, with Hit@1 values below 1% and MRR values below 0.02.

A hierarchy-aware evaluation further revealed that retrieval performance decreased with increasing taxonomic abstraction. For the best-performing model, Hit@1 was 3.25% at the *subclass* level and below 1% for *class*, *superclass*, and *kingdom*, indicating limited semantic alignment between substance names and higher-level ChemOnt categories. A weighted lenient MRR, designed to account for partial matches across the hierarchy, yielded a value of 0.099 for the MiniLM model, only marginally higher than the strict MRR.

Detailed performance metrics for all evaluated models are provided in Table 1.

Table 1. Performance of embedding models for ChemOnt classification.

Model	Hit@1	Hit@3	Hit@5	Hit@10	MRR
all-MiniLM-L6-v2	0.0575	0.0986	0.1234	0.1621	0.093
all-mpnet-base-v2	0.0093	0.0152	0.0184	0.0247	0.015
SciBERT	0.0040	0.0071	0.0102	0.0131	0.008
BioBERT	0.0047	0.0077	0.0099	0.0150	0.009

Scale of the Interoperability Problem For the more than 15,000 non-structure entries, full manual reconciliation would require on the order of hundreds of millions of pairwise comparisons. Under conservative assumptions, this corresponds to several thousand person-years of expert effort.

Figure 6 quantifies this workload across all regulatory sources, illustrating both the absolute scale and the quadratic growth that makes each additional non-structure entry disproportionately costly. Even moderate increases in the number of unresolved entries lead to a rapidly escalating curation burden.

Manual pairwise reconciliation at this scale is not operationally feasible. This result highlights a fundamental constraint: interoperability cannot be achieved through retrospective harmonisation of heterogeneous substance definitions.

Instead, the only scalable strategy is to address the problem at its source. By requiring structure-resolvable identifiers and ontology-based representations in regulatory definitions wherever chemically applicable, the number of entries requiring manual reconciliation can be structurally reduced rather than managed post hoc.

Knowledge Graph Integration Structure-defined substances can be successfully integrated across regulatory lists, chemical taxonomies (ChemOnt), and biological annotations (ChEBI) using InChIKey as a linking key. This demonstrates the feasibility of a linked-data approach when structure information is available.

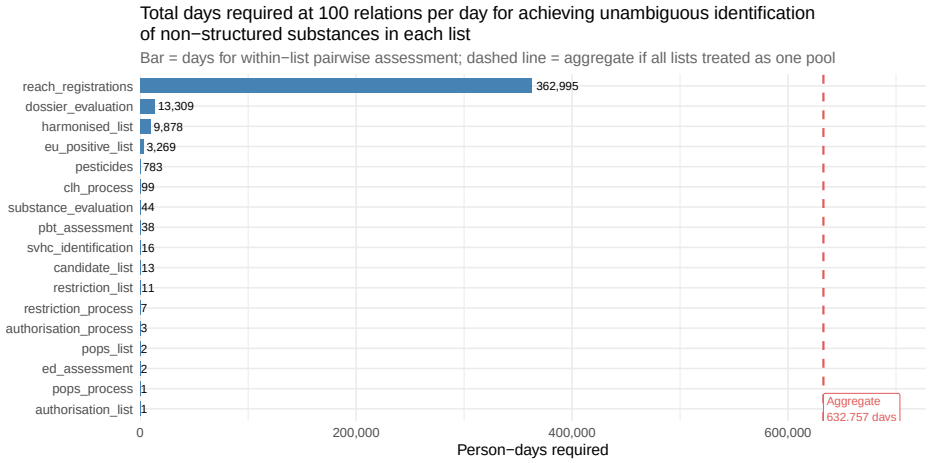

Figure 6. Estimated pairwise curation workload for 11,250 non-structure entries across the combined European regulatory dataset, broken down by source. The total corresponds to approximately 630,000 person–days (3,150 person-years) under an assumed throughput of 100 pair assessments per person-day. The quadratic growth of the workload illustrates that manual reconciliation is not scalable, reinforcing the need for structure-based and ontology-driven representations at the source.

Priority Patterns The integrated knowledge graph enables the combination of regulatory and biological annotations into a shared analytical space. Using the composite priority score defined in Section 3, we identify substances and chemical classes that are prominent across both regulatory presence and biological role annotations.

The bubble chart (Figure 7) visualises this priority space by plotting list_score against bio_score for all scored substances. Bubble size is proportional to the number of regulatory source memberships.

Substances located in the upper-right quadrant combine high regulatory coverage with multiple biological role annotations, reflecting their prominence across heterogeneous data sources. Examples include bisphenol A (BPA) and dibutyl phthalate, which appear consistently across regulatory lists and are associated with multiple biological roles in ChEBI.

Importantly, this analysis does not constitute a toxicological ranking, but demonstrates how interoperable data representations enable the integration of regulatory and ontology-based information within a single analytical framework.

At the class level, a small number of chemical classes accounts for a large fraction of regulated substances. This concentration suggests that class-based representations can provide a scalable bridge between individual substances and regulatory group definitions, supporting alignment across regulatory contexts.

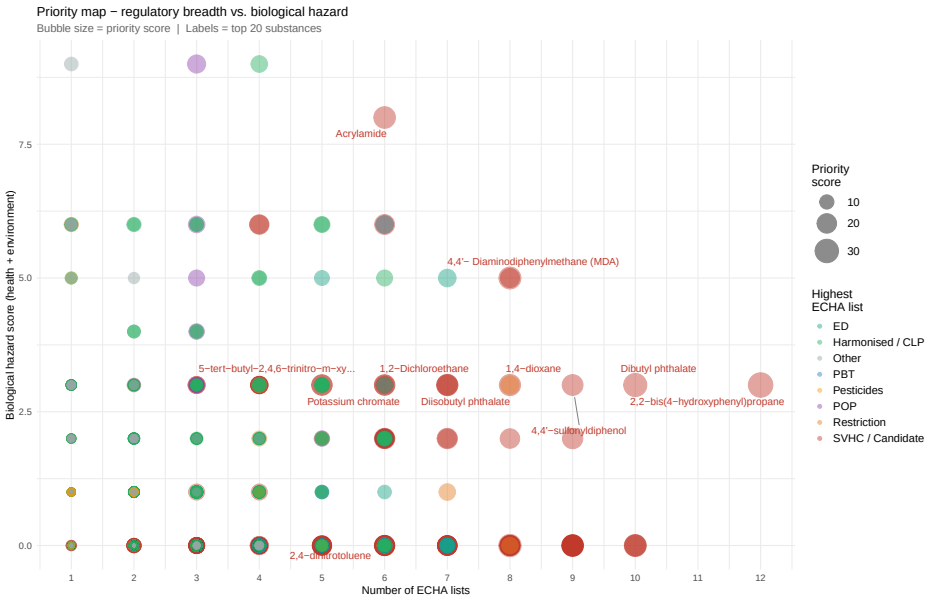

Figure 7. Bubble chart of composite priority space: regulatory list score versus biological hazard score for all scored substances. Bubble size is proportional to the number of regulatory source memberships. Substances in the upper-right quadrant combine broad regulatory coverage with documented biological hazard evidence.

5. Discussion

Chemical identity as a prerequisite for interoperability The results demonstrate that interoperability of regulatory substance data is fundamentally constrained by the absence of a shared and formalised notion of chemical identity. While structure-based identifiers such as InChIKey enable unambiguous linking across datasets, a substantial fraction of regulatory entries cannot be expressed in this form. Instead, regulatory "substances" frequently represent heterogeneous entities, including mixtures, substance groups, and analytical parameters.

This heterogeneity is not merely a data quality issue but reflects a structural mismatch between chemical reality and regulatory abstraction. As a consequence, semantic relationships between regulatory entries, such as equivalence, subset relations or partial overlap, remain implicit and undefined. Without explicit formalisation of these relationships, automated integration across regulatory datasets is inherently limited.

Limitations of identifier-based approaches The analysis confirms that commonly used identifiers in regulatory contexts, particularly CAS numbers, do not provide a reliable foundation for interoperability. Although widely adopted, they exhibit both one-to-many and many-to-one mappings with chemical structures, reflecting differences in stereochemistry, salt forms, and context-dependent definitions.

This finding challenges the widespread assumption in regulatory data integration that such identifiers can serve as globally unique references. In practice, relying on them introduces ambiguity rather than resolving it, undermining their suitability as a backbone for cross-dataset linking.

Limitations of semantic and AI-based approaches This study shows that modern embedding-based methods cannot compensate for the absence of structure-based identifiers in regulatory substance data. While they can group substance names into loosely coherent semantic clusters, they fail to recover true chemical identity or reliably map entries to ChemOnt classes, with only a negligible number of correct assignments.

These limitations are inherent to the data rather than the models: textual descriptions of mixtures, groups, and analytical entries lack the information required to infer molecular structure. As a result, semantic similarity largely reflects lexical overlap rather than chemical relatedness.

Overall, the results indicate that sentence embeddings cannot substitute for structure-resolvable identifiers. Reliable chemical classification and interoperability therefore require explicit structural representation, such as InChIKeys, at the source.

The combinatorial barrier to retrospective harmonisation The scale of the interoperability problem becomes evident when considering the effort required for manual reconciliation. For non-structure substances, establishing semantic relationships requires pairwise comparison, leading to quadratic growth in workload. The estimated effort of several thousand person-years demonstrates that retrospective harmonisation is not a viable strategy at scale.

This reinforces a key insight: interoperability cannot be achieved by correcting data after publication, but must be designed into regulatory data models from the outset.

Towards ontology-driven regulatory data infrastructures In contrast to regulatory datasets, cheminformatics infrastructures such as PubChem and ChEBI demonstrate that large-scale interoperability is achievable when chemical identity is explicitly represented and linked to formal ontologies. Structure-based identifiers, combined with hierarchical classifications and machine-readable knowledge graphs, enable consistent integration across domains.

The priority scoring example illustrates that once chemical identity is formalised and linked across datasets, higher-level aggregations combining regulatory and biological information become straightforward to compute.

The results of this study show that similar principles can be applied to regulatory data. Structure-defined substances can be successfully linked across datasets and enriched with chemical and biological annotations when represented using InChIKey and ontology-based classifications. Moreover, the observed concentration of substances within a limited number of chemical classes suggests that class-level representations can provide a scalable bridge between individual molecules and regulatory group definitions.

Implications for FAIR chemical data in Europe The findings have direct implications for the implementation of FAIR (Findable, Accessible, Interoperable, Reusable) principles in European chemical regulation. Current regulatory datasets fall short of interoperability primarily because they lack globally unique, persistent, and machine-readable identifiers grounded in chemical structure.

Achieving FAIR chemical substance data requires a shift from document-centric and identifier-fragmented approaches towards linked data infrastructures in which substances are represented as resolvable URIs, grounded in structure-based identifiers and connected through formal ontologies. Importantly, this shift is not purely technical: it requires alignment between data modelling practices and the legal definitions of substances in regulatory frameworks.

Limitations and scope This study focuses on datasets derived from ECHA and related European regulatory sources, and on the use of InChIKey as the primary structure-based identifier. While the results are consistent across the analysed datasets, further work is needed to assess the generality of these findings across other regulatory domains and jurisdictions.

Furthermore, not all regulatory concepts can or should be reduced to individual molecular structures. Mixtures, UVCBs, and analytical parameters are legitimate and necessary constructs in regulation. However, their current representation lacks formal semantics, limiting their integration with structure-based systems. Future work should explore hybrid models that explicitly represent these entities as sets, classes, or functions over well-defined chemical structures.

From data integration to regulatory design Taken together, the results show that the interoperability problem is not primarily a technical integration challenge, but a consequence of how substances are defined and represented in regulation. Efforts to improve interoperability must therefore extend beyond data engineering to include the design of regulatory concepts and data standards.

Embedding structure-based identifiers, ontology references, and machine-readable semantics directly into regulatory processes would significantly reduce ambiguity and enable scalable integration. Without such changes, the gap between regulatory data and scientific data infrastructures will continue to widen, limiting the effectiveness of data-driven environmental governance.

Concretely, this implies that regulatory data infrastructures should: (i) include structure-based identifiers (e.g. InChIKey) wherever applicable; (ii) represent substance groups and mixtures using explicit semantic models (e.g. set- or class-based representations); (iii) publish substance data as linked data with persistent URIs.

6. Conclusion

This paper demonstrates that the interoperability of chemical substance data within European regulatory frameworks is fundamentally constrained by the absence of a shared and formalised notion of chemical identity. An analysis of 18 regulatory datasets revealed that a substantial proportion of substances cannot be unambiguously linked to a defined molecular structure. Moreover, commonly used identifiers, such as CAS numbers, do not provide a sufficiently reliable basis for cross-dataset integration.

This limitation has significant implications for semantic interoperability. Regulatory substance definitions frequently lack formal semantics and structure-based identifiers, leaving relationships such as equivalence, subset relations, and partial overlap implicit. The issue is particularly pronounced for group- and parameter-based substances, which are more accurately conceptualised as sets of molecules or as functions defined over such sets. For example, regulatory groups like per- and polyfluoroalkyl substances (PFAS) represent collections of structurally related compounds, while analytical parameters such as “Sum PFOS + PFBA” correspond to measurements derived from subsets of these groups. In the absence of an explicit and formal representation of these set-theoretic relationships, the semantic connections between substances remain undefined, thereby complicating interoperability and hindering automated reasoning.

The challenge is further amplified by the scale and diversity of European regulatory data. Multiple overlapping lists of substances of concern have been developed within distinct legal and domain-specific contexts. Without a shared representation of chemical identity, these lists cannot be reliably linked or compared. Consequently, retrospective harmonisation, whether through manual curation or data-driven techniques such as semantic embeddings, faces intrinsic limitations and results in a combinatorial barrier to integration.

Our findings indicate that interoperability cannot be achieved solely through technical post-processing. Instead, it must be addressed at the level of data modelling and legislation. The adoption of structure-based identifiers, such as InChIKeys, in combination with formal ontologies and linked data principles, enables meaningful integration when implemented from the outset. These results underscore that the interoperability challenge is not primarily technical but stems from the way chemical substances are defined and represented within regulatory frameworks.

We therefore argue that achieving FAIR and interoperable chemical data in Europe requires a paradigm shift towards embedding structure-based identifiers, formal ontologies, and linked data principles directly into regulatory data models and legislative processes. Without such changes, the disconnect between regulatory datasets and scientific data infrastructures will persist, limiting the effectiveness of data-driven environmental and chemical policy.

Future research should focus on the development of hybrid semantic models capable of representing mixtures and substance groups, as well as on translating these conceptual advances into operational standards for regulatory data publication and exchange.

Data Availability

All data, code, and materials supporting this study are publicly available. The complete dataset, including the integrated knowledge graph, analysis scripts, and figures, is archived on Zenodo at `https://doi.org/10.5281/zenodo.xxxxxxx`.

The development version of the repository is available on GitHub: `https://github.com/gezever/A-Substance-Is-Not-Always-a-Substance/`.

Machine-readable metadata describing the dataset is provided using the DCAT standard. All generated data products are published as machine-readable RDF and described using DCAT metadata to support FAIR principles.

Supplementary Materials

Supplementary materials include:

- The integrated knowledge graph (RDF/Turtle)
- Processed datasets derived from 18 regulatory sources
- Analysis scripts and workflows
- All figures and visualisations
- DCAT metadata describing the dataset and its distributions
- The initial methods and results section that was to large for this paper

All materials are available via the Zenodo archive associated with this publication.

References

[1] European Parliament, Council of the European Union. Regulation (EC) No 1907/2006 concerning the Registration, Evaluation, Authorisation and Restriction of Chemicals (REACH). Official Journal of the European Union; 2006. L 396. Available from: http://data.europa.eu/eli/reg/2006/1907/2022-03-01.

[2] European Parliament, Council of the European Union. Regulation (EC) No 1272/2008 on classification, labelling and packaging of substances and mixtures (CLP). Official Journal of the European Union; 2008. L 353. Available from: http://data.europa.eu/eli/reg/2008/1272/oj.

[3] Weininger D. SMILES, a chemical language and information system. 1. Introduction to methodology and encoding rules. Journal of Chemical Information and Computer Sciences. 1988 Feb;28(1):31-6. Available from: https://doi.org/10.1021/ci00057a005.

[4] Koppenol WH, Corish J, García-Martínez J, Meija J, Reedijk J. How to name new chemical elements (IUPAC Recommendations 2016). 2016 2026;88(4):401-5. Available from: https://doi.org/10.1515/pac-2015-0802.

[5] Heller SR, McNaught A, Pletnev I, Stein S, Tchekhovskoi D. InChI, the IUPAC International Chemical Identifier. Journal of Cheminformatics. 2015 May;7(1):23. Available from: https://doi.org/10.1186/s13321-015-0068-4.

[6] Akhondi SA, Kors JA, Muresan S. Consistency of systematic chemical identifiers within and between small-molecule databases. Journal of Cheminformatics. 2012 Dec;4(1):35. Available from: https://doi.org/10.1186/1758-2946-4-35.

[7] Kim S, Thiessen PA, Bolton EE, Bryant SH. PUG-SOAP and PUG-REST: web services for programmatic access to chemical information in PubChem. Nucleic acids research. 2015 Jul;43(W1):W605-11. Available from: https://doi.org/10.1093/nar/gkv396.

[8] Degtyarenko K, de Matos P, Ennis M, Hastings J, Zbinden M, McNaught A, et al. ChEBI: a database and ontology for chemical entities of biological interest. Nucleic Acids Research. 2008 Jan;36(suppl_1):D344-50. Available from: https://doi.org/10.1093/nar/gkm791.

[9] Chemical Abstracts Service. CAS Registry; 2024. Accessed: 2026-04-08. https://www.cas.org.

[10] Dittmar PG, Stobaugh RE, Watson CE. The Chemical Abstracts Service Chemical Registry System. I. General Design. Journal of Chemical Information and Computer Sciences. 1976 May;16(2):111-21. Available from: https://doi.org/10.1021/ci60006a016.

[11] European Chemicals Agency (ECHA). About EC and List Numbers; n.d. Accessed April 8, 2026. Available from: https://echa.europa.eu/en/about-ec-and-list-numbers.

[12] Kutsarova SS, Yordanova DG, Karakolev YH, Stoeva S, Comber M, Hughes CB, et al. UVCB substances II: Development of an endpoint-nonspecific procedure for selection of computationally generated representative constituents. Environmental Toxicology and Chemistry. 2019 Mar;38(3):682-94. Available from: https://doi.org/10.1002/etc.4358.

[13] Wilkinson MD, Dumontier M, Aalbersberg IJ, Appleton G, Axton M, Baak A, et al. The FAIR Guiding Principles for scientific data management and stewardship. Scientific Data. 2016 Mar;3(1):160018. Available from: https://doi.org/10.1038/sdata.2016.18.

[14] Djoumbou Feunang Y, Eisner R, Knox C, Chepelev L, Hastings J, Owen G, et al. ClassyFire: automated chemical classification with a comprehensive, computable taxonomy. Journal of Cheminformatics. 2016 Nov;8(1):61. Available from: https://doi.org/10.1186/s13321-016-0174-y.

[15] Reimers N, Gurevych I. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics; 2019. p. 3982-92.

[16] Filtz E, Kirrane S, Polleres A. The linked legal data landscape: linking legal data across different countries. Artificial Intelligence and Law. 2021 Dec;29(4):485-539. Available from: https://doi.org/10.1007/s10506-021-09282-8.

[17] European Chemicals Agency (ECHA). Candidate List of Substances of Very High Concern; 2025. Available from: https://chem.echa.europa.eu/obligation-lists/candidateList.

[18] European Chemicals Agency (ECHA). Authorisation List; 2025. Available from: https://chem.echa.europa.eu/obligation-lists/authorisationList.

[19] European Chemicals Agency (ECHA). Restriction List; 2025. Available from: https://chem.echa.europa.eu/obligation-lists/restrictionList.

[20] European Chemicals Agency (ECHA). Persistent Organic Pollutants (POPs) List; 2025. Available from: https://chem.echa.europa.eu/obligation-lists/popsList.

[21] European Chemicals Agency (ECHA). EU Positive List; 2025. Available from: https://chem.echa.europa.eu/obligation-lists/euPositiveList.
[22] European Chemicals Agency (ECHA). Harmonised Classification and Labelling List (CLH); 2025. Available from: https://chem.echa.europa.eu/obligation-lists/clhList.
[23] European Chemicals Agency (ECHA). Restriction Process List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/restrictionProcess.
[24] European Chemicals Agency (ECHA). SVHC Identification List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/svhcIdentification.
[25] European Chemicals Agency (ECHA). Authorisation Process List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/authorisationProcess.
[26] European Chemicals Agency (ECHA). Dossier Evaluation List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/dossierEvaluation.
[27] European Chemicals Agency (ECHA). CLH Process List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/clhProcess.
[28] European Chemicals Agency (ECHA). Substance Evaluation List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/substanceEvaluation.
[29] European Chemicals Agency (ECHA). POPs Process List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/popsProcess.
[30] European Chemicals Agency (ECHA). PBT Assessment List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/pbtAssessment.
[31] European Chemicals Agency (ECHA). Endocrine Disruptor Assessment List; 2025. Available from: https://chem.echa.europa.eu/activity-lists/edAssessment.
[32] European Chemicals Agency (ECHA). REACH Registrations Database; 2025. Available from: https://chem.echa.europa.eu/.
[33] European Commission. EU Pesticides Database: Active Substances; 2025. Available from: https://ec.europa.eu/food/plant/pesticides/eu-pesticides-database/start/screen/active-substances.
[34] OSPAR Commission. OSPAR List of Chemicals for Priority Action; 2024. Available from: https://www.ospar.org/work-areas/hasec/hazardous-substances/priority-action.
[35] Vlaamse Overheid. Sommatie Stoffenlijst (Waterkwaliteitsmodel); 2025. Available from: https://datasets.omgeving.vlaanderen.be/be.vlaanderen.omgeving.data.id.distribution.codelijst-sommatie_stoffen.3.0.8.sommatie_stoffen.csv.
[36] Rousseeuw PJ. Silhouettes: A Graphical Aid to the Interpretation and Validation of Cluster Analysis. Journal of Computational and Applied Mathematics. 1987;20:53-65.
[37] Lloyd SP. Least Squares Quantization in PCM. IEEE Transactions on Information Theory. 1982;28(2):129-37.
[38] McInnes L, Healy J, Melville J. UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction. arXiv preprint arXiv:180203426. 2018. Available from: https://arxiv.org/abs/1802.03426.
[39] Manning CD, Raghavan P, Schütze H. Introduction to Information Retrieval. Cambridge: Cambridge University Press; 2008.
[40] Wang W, Wei F, Dong L, Bao H, Yang N, Zhou M. MiniLM: Deep Self-Attention Distillation for Task-Agnostic Compression of Pre-Trained Transformers. In: Advances in Neural Information Processing Systems. vol. 33; 2020. p. 5776-88. Available from: https://proceedings.neurips.cc/paper/2020/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html.
[41] Devlin J, Chang MW, Lee K, Toutanova K. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In: Burstein J, Doran C, Solorio T, editors. Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers). Minneapolis, Minnesota: Association for Computational Linguistics; 2019. p. 4171-86. Available from: https://aclanthology.org/N19-1423/.
[42] Miles A, Bechhofer S. SKOS Simple Knowledge Organization System Reference. World Wide Web Consortium; 2009. Available from: https://www.w3.org/TR/skos-reference/.
[43] Sanderson R, Ciccarese P, Young B. Web Annotation Data Model. World Wide Web Consortium; 2017. Available from: https://www.w3.org/TR/annotation-model/.
[44] European Commission. Chemicals Strategy for Sustainability Towards a Toxic-Free Environment. European Commission; 2020. COM(2020) 667 final. Available from: https://eur-lex.europa.eu/

legal-content/EN/TXT/?uri=CELEX:52020DC0667.

[45] European Parliament, Council of the European Union. Regulation (EU) 2019/1021 on persistent organic pollutants (POPs). Official Journal of the European Union; 2019. L 169. Available from: http://data.europa.eu/eli/reg/2019/1021/oj.

[46] Reif DM, Martin MT, Tan SW, Houck KA, Judson RS, Richard AM, et al. Endocrine Profiling and Prioritization of Environmental Chemicals Using ToxCast Data. Environmental Health Perspectives. 2010;118(12):1714-20.